

An Insufficient Introduction to Spark

Part 4: Relational Algebra and Table Processing

Riccardo Murri <riccardo.murri@gmail.com>

Relational Algebra

Relational DBs

All major general purpose DBMS's are based on the so-called relational data model.

This means that all data is stored in a number of tables (with named columns), such as:

usr	size	path
usr264	17	/scratch/ift/usr264/cp1.log
usr116	19362662400	/scratch/id/usr116/vkeller.tar
usr116	3379200	/scratch/id/usr116/test.tar
usr264	16	/scratch/ift/usr264/cp2.log
usr345	877366	/scratch/aim/usr345/bwa/bwa

For historical and mathematical reasons such tables are referred to as *relations*.

Relational data model

A **relational database** is a set of relations.

A **relation** is an (ordered) set of tuples.

A relation can be represented by listing all groups of related elements — the result is a “table”.

usr	size	path
usr264	17	/scratch/ift/usr264/cp1.log
usr116	19362662400	/scratch/id/usr116/vkeller.tar
usr116	3379200	/scratch/id/usr116/test.tar
usr264	16	/scratch/ift/usr264/cp2.log
usr345	877366	/scratch/aim/usr345/bwa/bwa

What is *relational algebra*?

Relational algebra, defined in its basic form by E. F. Codd in 1970, has:

- ▶ relations as atomic operands, and
- ▶ various operations on relations as operators (which will be detailed shortly).

Relational algebra is the basis of SQL and of Spark DataFrames.

Relational Algebra

DataFrame are basically tables: all the relational algebra operators that we already know can be applied.

`df.select(...)` return new DataFrame with only the specified set of columns (RA's "projection" operator).

`df.distinct()` return new DataFrame omitting duplicate rows.

`df.where(...)` return new DataFrame containing only rows that satisfy a condition (RA's "selection" operator).

`df.orderBy(...)` return new DataFrame sorted by the specified column(s).

Relational Algebra (2)

DataFrame are basically tables: all the relational algebra operators that we already know can be applied.

`df1.intersect(df2)` return new DataFrame containing only rows that are in *both* df1 and df2

`df.subtract(df2)` Return a new DataFrame containing only rows that are in df1 *but not* df2

`df.unionAll(df2)` Return a new DataFrame containing union of rows.

Relational Algebra (3)

`DataFrame` are basically tables: all the relational algebra operators that we already know can be applied.

`df1.join(df2, on=..., how=...)`

Perform a join of two `DataFrames`, return result as a new `DataFrame`.

The `how=` parameter is a string naming the type of JOIN operation: `'inner'`, `'outer'`, `'left_outer'`, `'right_outer'`.

The `on=` parameter is any of the following:

- ▶ `None` (default): perform a natural join
- ▶ column name or list of column names: perform an equi-join on the given columns
- ▶ column expression (or list thereof): perform a θ -join

Joins

Joins

From Wikipedia:

The JOIN operation combines columns from one or more tables in a relational database [...] by using values common to each.

There are 4 common types of JOIN: INNER, LEFT OUTER, RIGHT OUTER, and FULL OUTER.

Reference: For more details, see: [https://en.wikipedia.org/wiki/Join_\(SQL\)](https://en.wikipedia.org/wiki/Join_(SQL))

INNER JOIN

An `INNER JOIN` returns the set of tuples from the cross product of two tables that satisfy a certain predicate:

Table: Directors			Table: Prizes	
name	prizeId	⋈	id	prize
F. F. Coppola	11		27	Leone d'Oro
T. Kitano	27		11	Oscar

↓

Table: Result	
name	prize
F. F. Coppola	Oscar
T. Kitano	Leone d'Oro

```
directors.join(prizes,  
               on=(directors.prizeId == prizes.id),  
               how='inner')
```

Equi-joins

A join operation is called an *equi-join* if the selection predicate is a conjunction of equality comparisons.

In the case of an equi-join, the `on=...` argument can be omitted.

OUTER JOIN

In a OUTER JOIN, the result table retains each row, even if no other matching row exists.

- ▶ In a LEFT OUTER JOIN, every row from the **left** table is retained: the result is padded with NULL if no matching row from the *right* table is found.
- ▶ A RIGHT OUTER JOIN does the same with *left* and **right** reversed.
- ▶ In a FULL OUTER JOIN, rows from both sides are retained.

LEFT OUTER JOIN

Table: Directors			Table: Prizes	
name	prizeId		id	prize
F. F. Coppola	11	⋈	27	Leone d'Oro
T. Kitano	27		11	Oscar
Ed Wood	NULL			

↓

Table: Result	
name	prize
F. F. Coppola	Oscar
T. Kitano	Leone d'Oro
Ed Wood	NULL

```
directors.join(prizes,  
               on=(directors.prizeId == prizes.id),  
               how='left_outer')
```

SQL queries

SQL queries

It is possible to run SQL queries on DataFrame objects.

df.createTempView(*name*)

Allow queries on *df* as table *name*. Lifetime of the table is tied to SQLContext.

spark.sql(*query*)

Run *query* and return result as a new DataFrame. All registered DFs can be queried.

```
df.createTempView('data')  
df2 = spark.sql('SELECT_*_FROM_data;')
```